

```
In [4]: import matplotlib.pyplot as plt
import numpy as np
import random, math, copy
```

```
In [5]: # Set random seed for reproducibility
np.random.seed(42)

# Parameters
num_arms = 3

# Oracle theta of each arm
theta = np.array([0.7, 0.5, 0.4])
```

Problem 2: Implement classical bandit algorithms

1. The epsilon-greedy Algorithm

```
In [80]: def epsilon_greedy(epsilon, N, theta):
        """
        Implement the epsilon-greedy algorithm for a Bernoulli bandit problem.

        Parameters
        -----
        epsilon : float
            The probability of exploration.
        N : int
            Number of time steps.
        theta : array-like
            True success probabilities of each arm.

        Returns
        -----
        total_reward : float
            Total reward accumulated over N time steps.
        """
        Q = np.zeros(num_arms) # Estimated values for each arm
        counts = np.zeros(num_arms) # Count of how many times each arm is pulled
        total_reward = 0 # Total reward tracker

        # Initialization: Pull each arm once
        for arm in range(num_arms):
            reward = 1 if np.random.rand() < theta[arm] else 0
            counts[arm] = 1
            Q[arm] = reward
            total_reward += reward

        # Main loop: Epsilon-greedy exploration and exploitation
        for t in range(num_arms, N):
            if np.random.rand() < epsilon:
```

```

        # Exploration: choose a random arm
        arm = np.random.randint(num_arms)
    else:
        # Exploitation: choose the arm with the highest estimated value
        arm = np.argmax(Q)

    # Simulate pulling the chosen arm
    reward = 1 if np.random.rand() < theta[arm] else 0

    counts[arm] += 1
    Q[arm] += (1 / counts[arm]) * (reward - Q[arm])

    total_reward += reward

    return total_reward

```

2. The UCB (Upper Confidence Bound) Algorithm

```

In [81]: def ucb(c, N, theta):
    """
    Implement the UCB (Upper Confidence Bound) algorithm for a Bernoulli bandit problem.

    Parameters
    -----
    c : float
        Confidence level parameter for the UCB algorithm.
    N : int
        Number of time steps.
    theta : array-like
        True success probabilities of each arm.

    Returns
    -----
    rewards_history : array
        The rewards obtained at each time step.
    """

    Q = np.zeros(num_arms)
    counts = np.zeros(num_arms)
    total_reward = 0

    # Initialize by pulling each arm once
    for arm in range(num_arms):
        reward = 1 if np.random.rand() < theta[arm] else 0
        Q[arm] = reward
        counts[arm] = 1
        total_reward += reward

    for t in range(num_arms+1, N+1):
        # Avoid division by zero because each arm was pulled once
        ucb_values = Q + c * np.sqrt((2*np.log(t))/counts)
        arm = np.argmax(ucb_values)
        reward = 1 if np.random.rand() < theta[arm] else 0
        counts[arm] += 1

```

```

    Q[arm] += (1/counts[arm])*(reward - Q[arm])
    total_reward += reward
    return total_reward

```

3. TS (Thompson Sampling) Algorithm

```

In [82]: from scipy.stats import beta

def thompson_sampling(N, theta, alpha_init, beta_init):
    """
    Implement the Thompson Sampling (TS) algorithm for a Bernoulli bandit problem.

    Parameters
    -----
    N : int
        Number of time steps.
    theta : array-like
        True success probabilities of each arm.
    alpha_init : array-like
        Initial alpha parameters for the Beta distributions of each arm.
    beta_init : array-like
        Initial beta parameters for the Beta distributions of each arm.

    Returns
    -----
    rewards_history : array
        The rewards obtained at each time step.
    """
    alpha = alpha_init.copy()
    beta_ = beta_init.copy()
    total_reward = 0
    for t in range(N):
        sampled_thetas = [np.random.beta(alpha[j], beta_[j]) for j in range(num_arm)]
        arm = np.argmax(sampled_thetas)
        reward = 1 if np.random.rand() < theta[arm] else 0
        total_reward += reward
        alpha[arm] += reward
        beta_[arm] += 1 - reward
    return total_reward

```

Problem 3: Each experiment lasts for $N = 5000$ time slots, and we run each experiment 200 trials. Results are averaged over these 200 independent trials.

```

In [83]: # Parameters
N = 5000
num_trials = 200
epsilons = [0.1, 0.5, 0.9]
cs = [1, 5, 10]

```

```

# Two sets of prior parameters for TS
# Set 1: (1,1), (1,1), (1,1)
alpha_set_1 = np.array([1, 1, 1])
beta_set_1 = np.array([1, 1, 1])

# Set 2: (601,401), (401,601), (2,3)
alpha_set_2 = np.array([601, 401, 2])
beta_set_2 = np.array([401, 601, 3])

# True parameters of the arms (as per the oracle, but not known to the algorithm)
theta = np.array([0.7, 0.5, 0.4])

```

```

In [84]: # Epsilon-greedy
print("Epsilon-greedy results:")
for eps in epsilons:
    rewards = []
    for _ in range(num_trials):
        rewards.append(epsilon_greedy(eps, N, theta))
    mean_reward = np.mean(rewards)
    print(f"Epsilon = {eps}, Average total reward over {num_trials} trials: {mean_r

```

Epsilon-greedy results:

Epsilon = 0.1, Average total reward over 200 trials: 3408.44

Epsilon = 0.5, Average total reward over 200 trials: 3085.66

Epsilon = 0.9, Average total reward over 200 trials: 2748.215

```

In [85]: # UCB
print("\nUCB results:")
for c_val in cs:
    rewards = []
    for _ in range(num_trials):
        rewards.append(ucb(c_val, N, theta))
    mean_reward = np.mean(rewards)
    print(f"c = {c_val}, Average total reward over {num_trials} trials: {mean_rewar

```

UCB results:

c = 1, Average total reward over 200 trials: 3408.315

c = 5, Average total reward over 200 trials: 2979.74

c = 10, Average total reward over 200 trials: 2829.24

```

In [86]: # Thompson Sampling
print("\nThompson Sampling results:")
rewards_set_1 = []
for _ in range(num_trials):
    rewards_set_1.append(thompson_sampling(N, theta, alpha_set_1, beta_set_1))
mean_set_1 = np.mean(rewards_set_1)
print(f"Set 1 Priors (1,1),(1,1),(1,1), Average total reward: {mean_set_1}")

rewards_set_2 = []
for _ in range(num_trials):
    rewards_set_2.append(thompson_sampling(N, theta, alpha_set_2, beta_set_2))
mean_set_2 = np.mean(rewards_set_2)
print(f"Set 2 Priors (601,401),(401,601),(2,3), Average total reward: {mean_set_2}")

```

Thompson Sampling results:

Set 1 Priors (1,1),(1,1),(1,1), Average total reward: 3480.75

Set 2 Priors (601,401),(401,601),(2,3), Average total reward: 3492.41

Problem 4

4.1 Find the optimal results for each algorithm

```
In [94]: num_trials = 100

epsilon_values = np.arange(0, 0.5, 0.01)
average_rewards = []

for eps in epsilon_values:
    rewards = []
    for _ in range(num_trials):
        rewards.append(epsilon_greedy(eps, N, theta))
    average_rewards.append(np.mean(rewards))

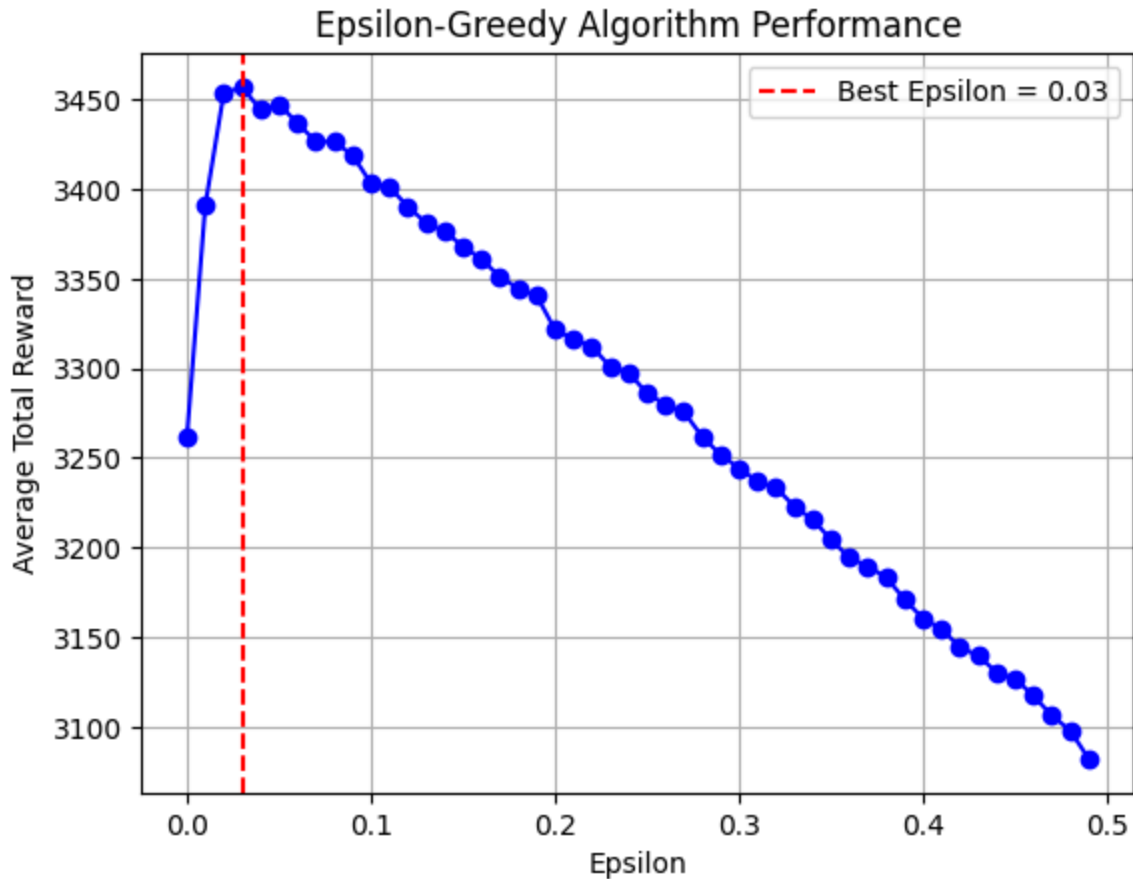
# Find the best epsilon
best_epsilon = epsilon_values[np.argmax(average_rewards)]
print(f"Best epsilon: {best_epsilon:.2f}")
print(f"Maximum average total reward: {np.max(average_rewards):.2f}")

# Plot the results
import matplotlib.pyplot as plt

plt.plot(epsilon_values, average_rewards, marker='o', linestyle='-', color = 'b')
plt.axvline(x=best_epsilon, color='r', linestyle='--', label=f'Best Epsilon = {best_epsilon:.2f}')
plt.xlabel('Epsilon')
plt.ylabel('Average Total Reward')
plt.title('Epsilon-Greedy Algorithm Performance')
plt.legend()
plt.grid(True)
plt.show()
```

Best epsilon: 0.03

Maximum average total reward: 3457.02



```
In [93]: c_values = np.arange(0, 5, 0.1)
         average_rewards = []

         # Run UCB for each value of c and compute the average reward over multiple trials
         for c in c_values:
             rewards = []
             for _ in range(num_trials):
                 total_reward = ucb(c, N, theta)
                 rewards.append(total_reward)
             average_rewards.append(np.mean(rewards))


         # Identify the best c
         best_c_index = np.argmax(average_rewards)
         best_c = c_values[best_c_index]
         best_average_reward = average_rewards[best_c_index]

         print(f"Best c value: {best_c:.2f}")
         print(f"Maximum average total reward: {best_average_reward:.2f}")

         # Plot the results
         plt.plot(c_values, average_rewards, marker='o', linestyle='-', color='b')
         plt.axvline(x=best_c, color='r', linestyle='--', label=f'Best c = {best_c:.2f}')
         plt.xlabel('c (Confidence Level Parameter)')
         plt.ylabel('Average Total Reward')
         plt.title('UCB Algorithm Performance vs c Parameter')
         plt.legend()
         plt.grid(True)
         plt.show()
```

Best c value: 0.40

Maximum average total reward: 3483.90



Problem 6

```
In [2]: import numpy as np
import random
import matplotlib.pyplot as plt
from itertools import product
import seaborn as sns
import pandas as pd

num_arms = 3
```

```
In [2]: # Initialize global variables for counts and estimated thetas
count = [0, 0, 0] # Corresponds to Arm 1, Arm 2, Arm 3
theta = [0.0, 0.0, 0.0] # Estimated thetas for Arm 1, Arm 2, Arm 3

def init_greedy():
    """
    Initializes the counts and estimated thetas for the greedy algorithm.
    """
    global count, theta
    count = [0, 0, 0] # Reset counts for Arms 1, 2, 3
    theta = [0.0, 0.0, 0.0] # Reset estimated thetas

def greedy_dependence(n, epsilon, initial_theta_oracle, p):
```

```

"""
Greedy algorithm with dependency in theta_oracle.

Parameters:
- n: Number of time steps
- epsilon: Exploration rate
- initial_theta_oracle: Initial probabilities for each arm [01, 02, 03]
- p: Probability adjustment parameter
"""

global count, theta
init_greedy() # Initialize counts and estimates
total_reward = 0 # Total actual rewards obtained

# Deep copy to avoid modifying the original initial_theta_oracle
current_theta = initial_theta_oracle.copy()

for t in range(n):
    prob = random.random() # Generate a random number in [0,1)

    if prob < epsilon:
        # Explore: choose a random arm from {0,1,2} corresponding to Arm 1, 2,
        arm = random.randint(0, 2)
    else:
        # Exploit: choose the arm with the highest estimated theta
        arm = np.argmax(theta)
        if theta[arm] == 0:
            # If all estimated thetas are 0, choose a random arm
            arm = random.randint(0, 2)

    # Simulate pulling the chosen arm: reward is 1 with probability current_theta
    r_i = np.random.binomial(1, current_theta[arm])

    # Accumulate the actual reward
    total_reward += r_i

    # Update counts and estimated thetas using incremental averaging
    count[arm] += 1
    theta[arm] += (r_i - theta[arm]) / count[arm]

    # Update theta_oracle based on the outcome
    if r_i == 1:
        # If reward obtained, decrease theta of pulled arm and increase others
        current_theta[arm] = max(current_theta[arm] - p, 0.0)
        for other_arm in range(3):
            if other_arm != arm:
                current_theta[other_arm] = min(current_theta[other_arm] + p / 2, 1.0)
    else:
        # If no reward, increase theta of pulled arm and decrease others
        current_theta[arm] = min(current_theta[arm] + p, 1.0)
        for other_arm in range(3):
            if other_arm != arm:
                current_theta[other_arm] = max(current_theta[other_arm] - p / 2, 0.0)

return total_reward

# Define the initial true reward probabilities (unknown to the algorithm)

```



```

initial_theta_oracle = [0.7, 0.5, 0.4] #  $\theta_1=0.7$ ,  $\theta_2=0.5$ ,  $\theta_3=0.4$ 

# Experiment Parameters
epsilon_values = np.arange(0, 1.02, 0.02) # Epsilon from 0 to 1 in steps of 0.02
repeat_time = 100 # Number of trials for each epsilon
N = 5000 # Number of time steps per trial
p = 0.005 # Probability adjustment parameter

rewards = np.zeros(len(epsilon_values)) # Average rewards for each epsilon

# Run experiments for each epsilon
for i, eps in enumerate(epsilon_values):
    for trial in range(repeat_time):
        # For each trial, reset the initial theta_oracle
        theta_oracle = initial_theta_oracle.copy()
        reward = greedy_dependence(N, eps, theta_oracle, p)
        rewards[i] += reward / repeat_time

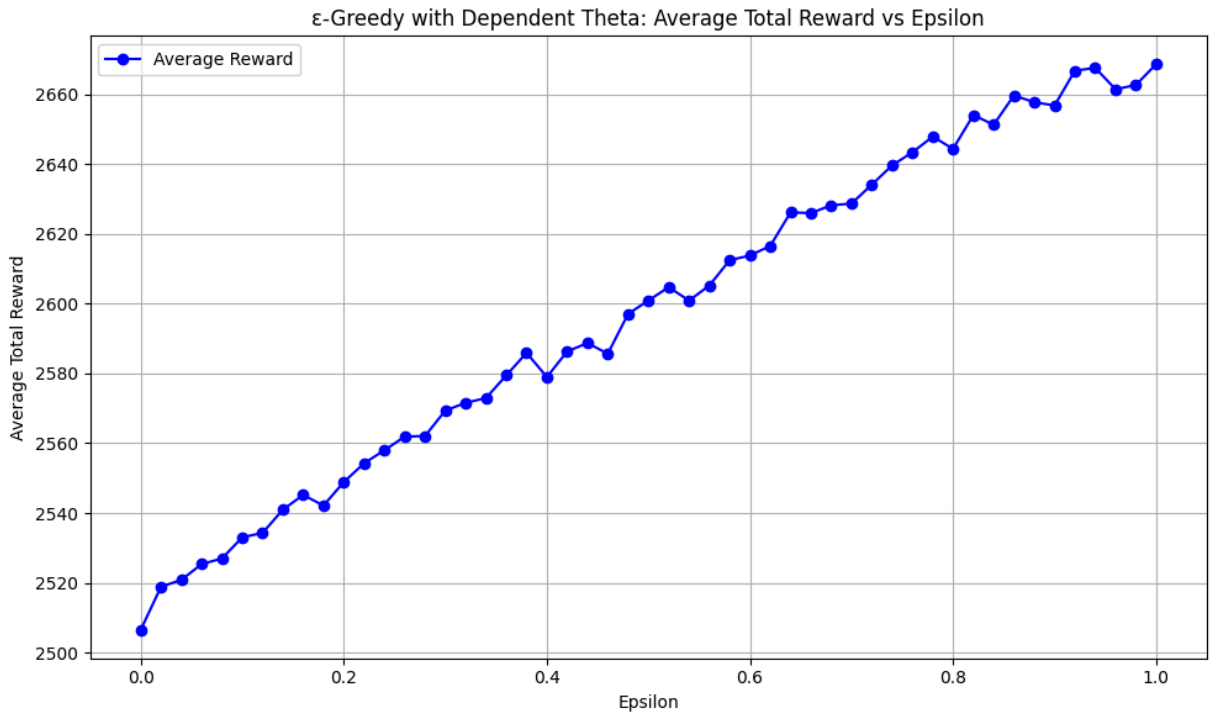
# Plot the results
plt.figure(figsize=(10, 6))

# Plot Average Total Reward vs. Epsilon
plt.plot(epsilon_values, rewards, marker='o', linestyle='-', color='blue', label='A')
plt.scatter(epsilon_values, rewards, color='red', s=10)
plt.xlabel('Epsilon')
plt.ylabel('Average Total Reward')
plt.title('ε-Greedy with Dependent Theta: Average Total Reward vs Epsilon')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

# Identify and print the best epsilon based on rewards
best_index = np.argmax(rewards)
best_epsilon = epsilon_values[best_index]
best_reward = rewards[best_index]
print(f"Best epsilon: {best_epsilon:.2f}")
print(f"Maximum average total reward: {best_reward:.2f}")

```



Best epsilon: 1.00

Maximum average total reward: 2668.57

```
In [3]: def ucb_dependence(c, N, initial_theta_oracle, p=0.005):
        """
        UCB algorithm with independent arms.

        Parameters:
        - c: Confidence parameter for UCB
        - N: Number of time steps
        - initial_theta_oracle: List of initial true reward probabilities for each arm
        - p: Probability adjustment parameter
        """
        num_arms = 3
        Q = np.zeros(num_arms)           # Estimated rewards for each arm
        counts = np.zeros(num_arms)      # Number of times each arm has been pulled
        total_reward = 0                  # Total accumulated reward

        # Deep copy to avoid modifying the original initial_theta_oracle
        theta_oracle = initial_theta_oracle.copy()

        # Initialize by pulling each arm once
        for arm in range(num_arms):
            reward = 1 if random.random() < theta_oracle[arm] else 0
            Q[arm] = reward
            counts[arm] = 1
            total_reward += reward

        # Update theta_oracle based on the outcome
        if reward == 1:
            # If reward obtained, decrease the probability of the pulled arm and in
            theta_oracle[arm] = max(theta_oracle[arm] - p, 0.0)
            for other_arm in range(num_arms):
                if other_arm != arm:
```

```

        theta_oracled[other_arm] = min(theta_oracled[other_arm] + p / 2
else:
    # If no reward, increase the probability of the pulled arm and decrease
    theta_oracled[arm] = min(theta_oracled[arm] + p, 1.0)
    for other_arm in range(num_arms):
        if other_arm != arm:
            theta_oracled[other_arm] = max(theta_oracled[other_arm] - p / 2

# Run UCB algorithm for remaining time steps
for t in range(num_arms, N):
    # Compute UCB values for each arm
    ucb_values = Q + c * np.sqrt((2 * np.log(t + 1)) / counts)
    arm = np.argmax(ucb_values)

    # Pull the selected arm and observe the reward
    reward = 1 if random.random() < theta_oracled[arm] else 0
    total_reward += reward

    # Update counts and estimated rewards
    counts[arm] += 1
    Q[arm] += (reward - Q[arm]) / counts[arm]

    # Update theta_oracled based on the outcome
    if reward == 1:
        # If reward obtained, decrease the probability of the pulled arm and in
        theta_oracled[arm] = max(theta_oracled[arm] - p, 0.0)
        for other_arm in range(num_arms):
            if other_arm != arm:
                theta_oracled[other_arm] = min(theta_oracled[other_arm] + p / 2
    else:
        # If no reward, increase the probability of the pulled arm and decrease
        theta_oracled[arm] = min(theta_oracled[arm] + p, 1.0)
        for other_arm in range(num_arms):
            if other_arm != arm:
                theta_oracled[other_arm] = max(theta_oracled[other_arm] - p / 2

return total_reward

# Define the initial true reward probabilities (unknown to the algorithm)
initial_theta_oracled = [0.7, 0.5, 0.4] # [θ1, θ2, θ3]

# Experiment Parameters
c_values = np.arange(0.0, 10.2, 0.2) # Confidence parameter c from 0 to 10 in step
repeat_time = 100 # Number of trials for each c
N = 5000 # Number of time steps per trial
p = 0.005 # Probability adjustment parameter

average_rewards = np.zeros(len(c_values)) # Average rewards for each c

# Run experiments for each c
for i, c in enumerate(c_values):
    for trial in range(repeat_time):
        # For each trial, reset the initial_theta_oracled
        theta_oracled = initial_theta_oracled.copy()
        reward = ucb_dependence(c, N, theta_oracled, p)
        average_rewards[i] += reward / repeat_time

```

```


# Plot the results
plt.figure(figsize=(12, 6))

# Plot Average Total Reward vs. Confidence Parameter c
plt.plot(c_values, average_rewards, marker='o', linestyle='-', color='blue', label=
plt.scatter(c_values, average_rewards, color='red', s=10)
plt.xlabel('Confidence Parameter c')
plt.ylabel('Average Total Reward')
plt.title('UCB: Average Total Reward vs Confidence Parameter c')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

# Identify and print the best c based on rewards
best_index = np.argmax(average_rewards)
best_c = c_values[best_index]
best_reward = average_rewards[best_index]
print(f"Best c: {best_c:.2f}")
print(f"Maximum average total reward: {best_reward:.2f}")

```



Best c: 8.60

Maximum average total reward: 2662.44

```

In [3]: def thompson_sampling_dependence(N, theta_oracle, alpha_init, beta_init, p=0.005):
    num_arms = 3
    alpha = alpha_init.copy()
    beta = beta_init.copy()
    total_reward = 0.0

    # Deep copy to avoid modifying the original theta_oracle
    theta_current = theta_oracle.copy()

    for t in range(N):
        # Sample theta from Beta distributions for each arm
        sampled_thetas = [np.random.beta(alpha[j], beta[j]) for j in range(num_arms)

```

```

# Select the arm with the highest sampled theta
arm = np.argmax(sampled_thetas)

# Simulate pulling the selected arm: reward is 1 with probability theta_cur
reward = 1 if np.random.random() < theta_current[arm] else 0
total_reward += reward

# Update the Beta distribution parameters for the selected arm
alpha[arm] += reward
beta[arm] += (1 - reward)

# Update theta_oracle based on the outcome
if reward == 1:
    # If reward obtained, decrease theta of pulled arm and increase others
    theta_current[arm] = max(theta_current[arm] - p, 0.0)
    for other_arm in range(num_arms):
        if other_arm != arm:
            theta_current[other_arm] = min(theta_current[other_arm] + p / 2, 1.0)
else:
    # If no reward, increase theta of pulled arm and decrease others
    theta_current[arm] = min(theta_current[arm] + p, 1.0)
    for other_arm in range(num_arms):
        if other_arm != arm:
            theta_current[other_arm] = max(theta_current[other_arm] - p / 2, 0.0)

return total_reward

```

```

In [4]: def dependency_aware_thompson_sampling(N, theta, alpha_init, beta_init, p=0.005, epsilon=0.01):
    alpha = alpha_init.copy()
    beta = beta_init.copy()
    K = len(alpha)
    theta_current = theta.copy()
    total_reward = 0

    for t in range(N):
        # Epsilon-greedy: with prob epsilon, pick a random arm
        if np.random.rand() < epsilon:
            chosen_arm = np.random.choice(K)
        else:
            # Otherwise, Thompson sample from each arm's Beta posterior
            samples = [np.random.beta(alpha[i], beta[i]) for i in range(K)]
            chosen_arm = np.argmax(samples)

        # Observe reward from environment
        reward = (np.random.rand() < theta_current[chosen_arm])
        total_reward += reward

        # --- Update Beta posterior for chosen arm and other arms ---
        if reward:
            # Chosen arm gets a standard Beta update
            alpha[chosen_arm] += 1

            # If gamma > 0, nudge alpha of the other arms
            for other_arm in range(K):
                if other_arm != chosen_arm:

```

```

        alpha[other_arm] += gamma
    else:
        # Chosen arm gets a standard Beta update
        beta[chosen_arm] += 1

        # If gamma > 0, nudge beta of the other arms
        for other_arm in range(K):
            if other_arm != chosen_arm:
                beta[other_arm] += gamma

    # environment update
    if reward:
        theta_current[chosen_arm] = max(theta_current[chosen_arm] - p, 0.0)
        for other_arm in range(K):
            if other_arm != chosen_arm:
                theta_current[other_arm] = min(theta_current[other_arm] + p/2,
    else:
        theta_current[chosen_arm] = min(theta_current[chosen_arm] + p, 1.0)
        for other_arm in range(K):
            if other_arm != chosen_arm:
                theta_current[other_arm] = max(theta_current[other_arm] - p/2,

    return total_reward

```

```

In [6]: # Define the true reward probabilities (independent arms)
theta1_true = 0.5
theta2_true = 0.4
theta3_true = 0.7
theta = [theta3_true, theta1_true, theta2_true] # [θ1, θ2, θ3] = [0.7, 0.5, 0.4]

# ----- Parameter Ranges for Three Arms -----
alpha1_values = [450, 600]
beta1_values = [1, 16]
alpha2_values = [300, 450]
beta2_values = [1, 16]
alpha3_values = [150, 300, 450]
beta3_values = [16, 31]

N = 5000
repeat_time = 50

# ----- Generate ALL Parameter Combinations for Three Arms -----
parameter_combinations = list(product(
    alpha1_values, alpha2_values, alpha3_values,
    beta1_values, beta2_values, beta3_values
))

# ----- Initialize Result Lists for Both Algorithms -----
results_ts_independent = [] # For thompson_sampling_independent
results_da_ts_independent = [] # For dependency_aware_thompson_sampling (now indep

# ----- Running Both Algorithms Across ALL Parameter Combinations -----
for idx, (alpha1_val, alpha2_val, alpha3_val, beta1_val, beta2_val, beta3_val) in e
    alpha_init_ts = [alpha1_val, alpha2_val, alpha3_val]
    beta_init_ts = [beta1_val, beta2_val, beta3_val]

```

```

alpha_init_da = [alpha1_val, alpha2_val, alpha3_val]
beta_init_da = [beta1_val, beta2_val, beta3_val]

# ----- Run Trials for Thompson Sampling Independent -----
total_reward_ts = 0.0
for _ in range(repeat_time):
    theta_oracle = [0.7, 0.5, 0.4] # [θ1, θ2, θ3]
    reward = thompson_sampling_dependence(N, theta_oracle, alpha_init_ts, beta_init_ts)
    total_reward_ts += reward / repeat_time

# ----- Run Trials for Dependency-Aware Thompson Sampling -----
epsilon_values = [1e-2, 7e-3, 5e-3, 3e-3, 1e-3]
gamma_values = [1e-6, 1e-5, 1e-4, 1e-3, 1e-2]

best_reward_da = -np.inf
best_epsilon = None
best_gamma = None

for epsilon in epsilon_values:
    for gamma in gamma_values:
        total_reward_da = 0.0
        for _ in range(repeat_time):
            # Reset theta_oracle for each trial
            theta_oracle = [theta3_true, theta1_true, theta2_true] # [θ1, θ2, θ3]
            # Run DA-TS with current epsilon and gamma
            reward = dependency_aware_thompson_sampling(
                N, theta_oracle, alpha_init_da, beta_init_da,
                p=0.005, epsilon=epsilon, gamma=gamma
            )
            total_reward_da += reward / repeat_time

        # Check if this (epsilon, gamma) pair yields a better reward
        if total_reward_da > best_reward_da:
            best_reward_da = total_reward_da
            best_epsilon = epsilon
            best_gamma = gamma

# ----- Store Results -----
results_ts_independent.append({
    'Alpha1': alpha1_val,
    'Alpha2': alpha2_val,
    'Alpha3': alpha3_val,
    'Beta1': beta1_val,
    'Beta2': beta2_val,
    'Beta3': beta3_val,
    'Avg Reward TS': total_reward_ts
})
results_da_ts_independent.append({
    'Alpha1': alpha1_val,
    'Alpha2': alpha2_val,
    'Alpha3': alpha3_val,
    'Beta1': beta1_val,
    'Beta2': beta2_val,
    'Beta3': beta3_val,
    'Best Epsilon DA-TS': best_epsilon,
    'Best Gamma DA-TS': best_gamma,
})

```

```

        'Avg Reward DA-TS': best_reward_da
    })

# ----- Convert Results to DataFrames -----
df_ts = pd.DataFrame(results_ts_independent)
df_da = pd.DataFrame(results_da_ts_independent)

# ----- Merge DataFrames for Easier Comparison -----
df_combined = pd.merge(df_ts, df_da, on=['Alpha1', 'Alpha2', 'Alpha3', 'Beta1', 'Be

# ----- Find Best Outcomes for Each Method -----

best_avg_ts = df_ts['Avg Reward TS'].max()
best_avg_da = df_da['Avg Reward DA-TS'].max()

# ----- Enhanced Printing -----

print("==== Best Average Rewards =====")
print(f"Thompson Sampling Independent: {best_avg_ts:.2f}")
print(f"Dependency-Aware Thompson Sampling Independent: {best_avg_da:.2f}")

# ----- List Top 5 Parameter Combinations for Each Method -----

print("\n==== Top 5 Parameter Combinations for Thompson Sampling Independent =====")
top5_ts = df_ts.sort_values(by='Avg Reward TS', ascending=False).head(5)
print(top5_ts.to_string(index=False))

print("\n==== Top 5 Parameter Combinations for Dependency-Aware Thompson Sampling")
top5_da = df_da.sort_values(by='Avg Reward DA-TS', ascending=False).head(5)
print(top5_da.to_string(index=False))

# ----- Plotting -----

# Set the style for seaborn
sns.set(style="whitegrid")

# 1. Bar Plot of Best Average Rewards
greedy_dependence_avg = 2668.57
ucb_dependence_avg = 2662.44
plt.figure(figsize=(10, 6))
methods = ['Greedy Dependence', 'UCB Dependence', 'TS Independent', 'Dependency-Awa
avg_rewards = [greedy_dependence_avg, ucb_dependence_avg, best_avg_ts, best_avg_da]
# Fix deprecated palette usage in barplot
sns.barplot(x=methods,
            y=avg_rewards,
            hue=methods, # Assign x to hue
            legend=False, # Hide redundant Legend
            palette="viridis")
plt.ylabel('Average Reward')
plt.title('Best Average Rewards for Each Method')
plt.ylim(0, max(avg_rewards)*1.1)
for i, v in enumerate(avg_rewards):
    plt.text(i, v + max(avg_rewards)*0.01, f"{v:.2f}", ha='center', fontweight='bol
plt.tight_layout()
plt.show()

```



```

# 2. Scatter Plot Comparing Both Methods with Reference Line
plt.figure(figsize=(12, 8))
scatter = sns.scatterplot(
    data=df_combined,
    x='Avg Reward TS',
    y='Avg Reward DA-TS',
    hue='Alpha1',
    size='Beta1',
    palette='deep',
    alpha=0.7
)
# Add reference line y = x
max_val = max(df_combined['Avg Reward TS'].max(), df_combined['Avg Reward DA-TS'].max())
plt.plot([2600, max_val], [2600, max_val], 'r--', label='y = x')
plt.xlabel('Average Reward TS Independent')
plt.ylabel('Average Reward DA-TS Independent')
plt.title('Comparison of Average Rewards: TS vs DA-TS')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

# 3. Optional: Save Plots
plt.savefig('pics/best_average_rewards.png')
plt.savefig('pics/comparison_scatter.png')

```

===== Best Average Rewards =====

Thompson Sampling Independent: 2667.74


Dependency-Aware Thompson Sampling Independent: 2673.80

===== Top 5 Parameter Combinations for Thompson Sampling Independent =====

Alpha1	Alpha2	Alpha3	Beta1	Beta2	Beta3	Avg Reward TS
600	300	450	1	1	31	2667.74
450	450	300	1	16	16	2664.94
450	450	450	1	1	31	2663.02
450	450	450	1	16	31	2663.00
450	300	450	1	1	31	2662.90

===== Top 5 Parameter Combinations for Dependency-Aware Thompson Sampling Independent =====

Alpha1	Alpha2	Alpha3	Beta1	Beta2	Beta3	Best Epsilon DA-TS	Best Gamma DA-TS
600	450	450	16	16	31	0.001	0.000001
450	450	450	1	1	31	0.007	0.000001
600	450	450	1	1	31	0.007	0.000001
450	300	300	1	1	31	0.010	0.010000
600	450	300	16	1	16	0.007	0.000100



<Figure size 640x480 with 0 Axes>

Part II

Problem 1: One intuitive policy suggests that in each time slot we should pull the arm for which the current expected value of θ_i is the largest. This policy behaves very good in most cases. Please design simulations to check the behavior of this policy

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import beta

np.random.seed(42)
```

```
In [9]: # Simulation parameters
true_theta = [0.7, 0.5]           # True success probabilities for arm 1 and arm 2
alpha_prior = [1, 1]              # Prior alpha parameters for Beta distributions
beta_prior = [1, 1]               # Prior beta parameters for Beta distributions
gamma_values = np.linspace(0.95, 1.0, 50) # Gamma values from 0.5 to 1.0 in increments of 0.05
time_steps = 5000                  # Number of pulls per trial
repeat_time = 50                   # Number of trials per gamma

# Initialize array to store average total rewards for each gamma
average_total_rewards = []

# Iterate over each gamma value
for gamma in gamma_values:
    total_rewards = [] # To store total rewards for each trial

    # Repeat the trial 'repeat_time' times for averaging
    for trial in range(repeat_time):
        # Initialize Beta parameters for each arm
        alpha = alpha_prior.copy()
        beta_params = beta_prior.copy()

        cumulative_reward = 0 # Total reward for this trial

        for t in range(1, time_steps + 1):
            # Calculate expected theta for each arm
            expected_theta = [alpha[i] / (alpha[i] + beta_params[i]) for i in range(2)]

            # Select the arm with the highest expected theta
            chosen_arm = np.argmax(expected_theta)

            # Simulate a pull: success with probability true_theta[chosen_arm]
            success = np.random.rand() < true_theta[chosen_arm]

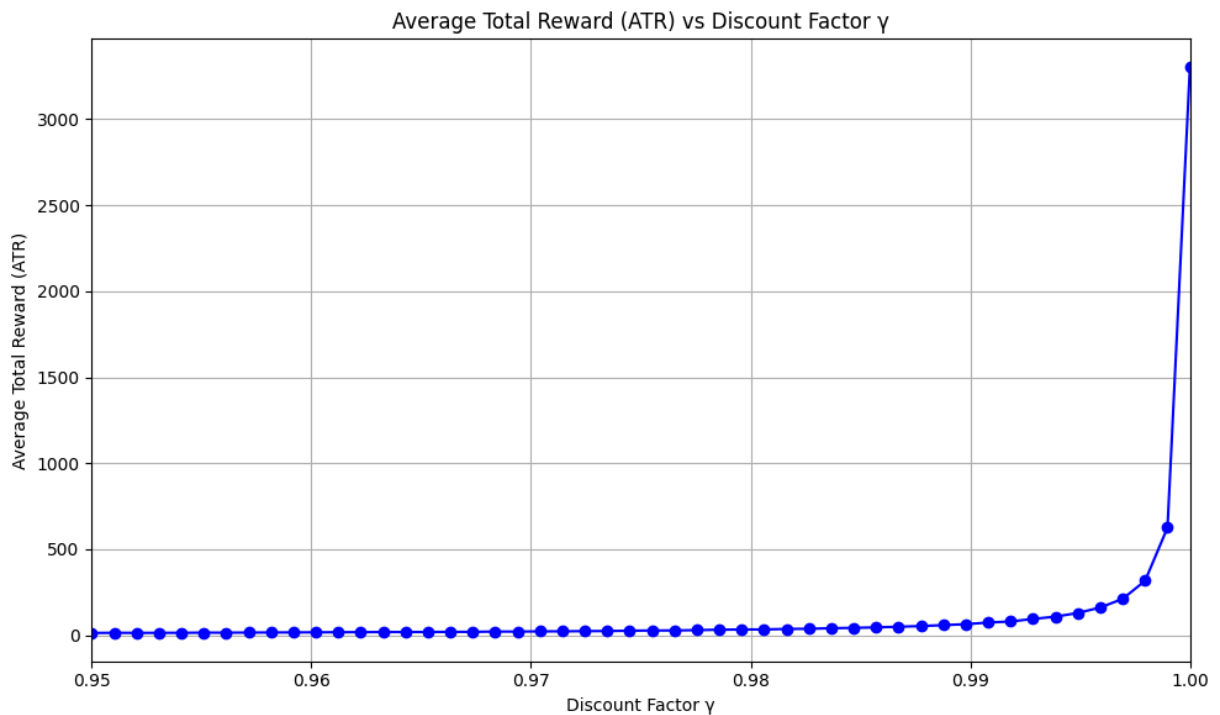
            # Update Beta posterior
            if success:
                alpha[chosen_arm] += 1
                reward = gamma**(t-1)
            else:
                beta_params[chosen_arm] += 1
                reward = 0

            # Update cumulative reward
            cumulative_reward += reward

        total_rewards.append(cumulative_reward)

    # Calculate average total reward for this gamma
    avg_reward = np.mean(total_rewards)
    average_total_rewards.append(avg_reward)
```

```
# Visualization: ATR vs Gamma
plt.figure(figsize=(10, 6))
plt.plot(gamma_values, average_total_rewards, marker='o', linestyle='-', color='blue')
plt.title('Average Total Reward (ATR) vs Discount Factor  $\gamma$ ')
plt.xlabel('Discount Factor  $\gamma$ ')
plt.ylabel('Average Total Reward (ATR)')
plt.grid(True)
plt.xlim(0.95, 1.0)
plt.tight_layout()
plt.show()
```



```
In [10]: # Identify the best arm (with the highest true_theta)
best_arm = np.argmax(true_theta)
theta_best = true_theta[best_arm]

# Initialize lists to store results
gamma_list = []
gap_list = []

# Iterate over each gamma value
for gamma in gamma_values:
    total_rewards = [] # To store total rewards for each trial

    # Repeat the trial 'repeat_time' times for averaging
    for trial in range(repeat_time):
        # Initialize Beta parameters for each arm
        alpha = alpha_prior.copy()
        beta_params = beta_prior.copy()

        cumulative_reward = 0.0 # Total reward for this trial

        for t in range(1, time_steps + 1):
            # Calculate expected theta for each arm using current Beta posterior
```

```

    expected_theta = [alpha[i] / (alpha[i] + beta_params[i]) for i in range

# Select the arm with the highest expected theta
chosen_arm = np.argmax(expected_theta)

# Simulate a pull: success with probability true_theta[chosen_arm]
success = np.random.rand() < true_theta[chosen_arm]

# Update Beta posterior based on the outcome
if success:
    alpha[chosen_arm] += 1
    reward = gamma**(t-1)
else:
    beta_params[chosen_arm] += 1
    reward = 0.0

# Accumulate the reward
cumulative_reward += reward

total_rewards.append(cumulative_reward)

# Calculate average total reward for this gamma
avg_reward = np.mean(total_rewards)

# Compute theoretical optimal reward
if gamma < 1.0:
    # Geometric series sum: theta_best * (1 - gamma^time_steps) / (1 - gamma)
    optimal_reward = theta_best * (1 - gamma**time_steps) / (1 - gamma)
else:
    # Handle the case when gamma = 1.0
    optimal_reward = theta_best * time_steps


# Compute the gap between optimal reward and algorithm's average reward
gap = optimal_reward - avg_reward

# Store the results
gamma_list.append(gamma)
gap_list.append(gap)

# Convert lists to numpy arrays for easier handling
gamma_array = np.array(gamma_list)
gap_array = np.array(gap_list)

# Visualization: Gap vs Gamma
plt.figure(figsize=(12, 6))
plt.plot(gamma_array, gap_array, marker='o', linestyle='--', color='red')
plt.title('Gap Between Optimal Reward and Algorithm\'s ATR vs Discount Factor  $\gamma$ ', f
plt.xlabel('Discount Factor  $\gamma$ ', fontsize=12)
plt.ylabel('Gap (Optimal - ATR)', fontsize=12)
plt.grid(True)
plt.xlim(0.95, 1.0)
plt.tight_layout()
plt.show()

```



Problem 2

However, such intuitive policy is unfortunately not optimal. Please provide an example to show why such policy is not optimal.

```
In [2]: # Simulation parameters
true_theta = [0.3, 0.6] # True success probabilities for arm 1 and arm 2
alpha_prior = [1, 1] # Prior alpha parameters for Beta distributions
beta_prior = [1, 1] # Prior beta parameters for Beta distributions
gamma_values = np.linspace(0.99, 1.00, 100) # Gamma values
time_steps = 5000 # Number of pulls per trial
repeat_time = 50 # Number of trials per gamma

# Identify the best arm (with the highest true_theta)
best_arm = np.argmax(true_theta)
theta_best = true_theta[best_arm]

# Initialize lists to store results
gamma_list = []
gap_list = []

# Iterate over each gamma value
for gamma in gamma_values:
    total_rewards = [] # To store total rewards for each trial

    # Repeat the trial 'repeat_time' times for averaging
    for trial in range(repeat_time):
        # Initialize Beta parameters for each arm
        alpha = alpha_prior.copy()
        beta_params = beta_prior.copy()

        cumulative_reward = 0.0 # Total reward for this trial

        for t in range(1, time_steps + 1):
            # Calculate expected theta for each arm using current Beta posterior
            expected_theta = [alpha[i] / (alpha[i] + beta_params[i]) for i in range
```

```

# Select the arm with the highest expected theta
chosen_arm = np.argmax(expected_theta)

# Simulate a pull: success with probability true_theta[chosen_arm]
success = np.random.rand() < true_theta[chosen_arm]

# Update Beta posterior based on the outcome
if success:
    alpha[chosen_arm] += 1
    reward = gamma**(t-1)
else:
    beta_params[chosen_arm] += 1
    reward = 0.0

# Accumulate the reward
cumulative_reward += reward

total_rewards.append(cumulative_reward)

# Calculate average total reward for this gamma
avg_reward = np.mean(total_rewards)

# Compute theoretical optimal reward
if gamma < 1.0:
    # Geometric series sum: theta_best * (1 - gamma^time_steps) / (1 - gamma)
    optimal_reward = theta_best * (1 - gamma**time_steps) / (1 - gamma)
else:
    # Handle the case when gamma = 1.0
    optimal_reward = theta_best * time_steps


# Compute the gap between optimal reward and algorithm's average reward
gap = optimal_reward - avg_reward

# Store the results
gamma_list.append(gamma)
gap_list.append(gap)

# Convert lists to numpy arrays for easier handling
gamma_array = np.array(gamma_list)
gap_array = np.array(gap_list)

# Visualization: Gap vs Gamma
plt.figure(figsize=(12, 6))
plt.plot(gamma_array, gap_array, marker='o', linestyle='--', color='red')
plt.title('Gap Between Optimal Reward and Algorithm\'s ATR vs Discount Factor  $\gamma$ ', f
plt.xlabel('Discount Factor  $\gamma$ ', fontsize=12)
plt.ylabel('Gap (Optimal - ATR)', fontsize=12)
plt.grid(True)
plt.xlim(0.99, 1.0)
plt.tight_layout()
plt.show()

```



Compared with TS

```
In [6]: # Simulation parameters
true_theta = [0.3, 0.6] # True success probabilities for arm 1 and arm 2
alpha_prior = [1, 1] # Prior alpha parameters for Beta distributions
beta_prior = [1, 1] # Prior beta parameters for Beta distributions
gamma_values = np.linspace(0.98, 1.00, 100) # Gamma values
time_steps = 5000 # Number of pulls per trial
repeat_time = 50 # Number of trials per gamma

# -----
# Helper function: sample Bernoulli reward from an arm
# -----
def draw_reward(arm_idx):
    """Simulate pulling arm_idx and return reward (1 or 0)."""
    return 1 if (np.random.rand() < true_theta[arm_idx]) else 0
# -----
# Naive Strategy: Always pick arm with highest posterior mean
# -----
def run_naive_strategy(gamma, alpha0, beta0):
    """
    Runs the naive strategy for 'time_steps' pulls with discount factor gamma.
    alpha0, beta0 are the prior parameters for each arm (list of length 2).
    Returns the total discounted reward.
    """
    # Initialize alpha, beta
    alpha = np.array(alpha0, dtype=float)
    beta = np.array(beta0, dtype=float)

    total_discounted_reward = 0.0
    discount_power = 0 # exponent for gamma^(t-1)

    for t in range(1, time_steps + 1):
        # Posterior means for each arm
        posterior_means = alpha / (alpha + beta)
```



```

    # Choose the arm with the highest posterior mean
    chosen_arm = np.argmax(posterior_means)

    # Draw a Bernoulli reward
    reward = draw_reward(chosen_arm)

    # Update posterior
    alpha[chosen_arm] += reward
    beta[chosen_arm] += (1 - reward)

    # Accumulate discounted reward
    total_discounted_reward += (gamma ** discount_power) * reward
    discount_power += 1

    return total_discounted_reward

# -----
# Thompson Sampling Strategy
# -----
def thompson_sampling(gamma, alpha0, beta0):
    """
    Runs Thompson Sampling for 'time_steps' pulls with discount factor gamma.
    alpha0, beta0 are the prior parameters for each arm.
    Returns the total discounted reward.
    """
    alpha = np.array(alpha0, dtype=float)
    beta = np.array(beta0, dtype=float)

    total_discounted_reward = 0.0
    discount_power = 0

    for t in range(1, time_steps + 1):
        # Sample theta-hat from current posterior for each arm
        sampled_thetas = np.random.beta(alpha, beta)

        # Choose the arm that maximizes the sampled theta
        chosen_arm = np.argmax(sampled_thetas)

        # Draw reward
        reward = draw_reward(chosen_arm)

        # Update posterior
        alpha[chosen_arm] += reward
        beta[chosen_arm] += (1 - reward)

        # Accumulate discounted reward
        total_discounted_reward += (gamma ** discount_power) * reward
        discount_power += 1

    return total_discounted_reward

# -----
# Main Experiment Loop
# -----
gap_means = []

for gamma in gamma_values:

```


```
gap_results = []

for _ in range(repeat_time):
    # Run Naive
    naive_reward = run_naive_strategy(
        gamma,
        alpha_prior,
        beta_prior
    )

    # Run Thompson Sampling
    ts_reward = thompson_sampling(
        gamma,
        alpha_prior,
        beta_prior
    )
    gap_results.append(ts_reward - naive_reward)

gap_means.append(np.mean(gap_results))

# -----
# Plotting Results
# -----
plt.figure(figsize=(8, 5))
plt.plot(gamma_values, gap_means, 'b-', label='Gaps between Naive and Thompson Samp
plt.xlabel('Discount Factor (gamma)')
plt.ylabel('Average Discounted Reward')
plt.title('Comparing Naive vs Thompson Sampling')
plt.legend()
plt.grid(True)
plt.show()
```



Problem 5

Find the optimal policy (approximately).

```
In [2]: import numpy as np
import matplotlib.pyplot as plt

# -----
# Dynamic Programming Function
# -----
def solve_2armed_bandit_dp(M=10, gamma=0.95, tol=1e-8, max_iter=50):
    """
    Solve the 2-armed Beta-Bernoulli bandit using 4D dynamic programming.

    Arguments:
    -----
    M : int
        Truncation level for alpha_i, beta_i.
    gamma : float
        Discount factor in (0,1).
    tol : float
        Convergence tolerance for the value iteration.
    max_iter : int
        Maximum number of iterations to run.

    Returns:
    -----
    R : 4D numpy array, shape (M+1, M+1, M+1, M+1)
```

```

    The approximate value function.
    policy : 4D numpy array of 0 or 1
    Optimal action: 0 for arm1, 1 for arm2.
    """
    # Initialize value function and policy arrays
    R = np.zeros((M+1, M+1, M+1, M+1), dtype=np.float64)
    policy = np.zeros((M+1, M+1, M+1, M+1), dtype=int)

    def clamp(x):
        return min(x, M)

    for it in range(max_iter):
        delta = 0.0 # Maximum change in this iteration

        # Iterate over all possible states
        for alpha1 in range(1, M+1):
            for beta1 in range(1, M+1):
                for alpha2 in range(1, M+1):
                    for beta2 in range(1, M+1):
                        # Compute expected reward for choosing arm 1
                        p1 = alpha1 / (alpha1 + beta1)
                        R_success_1 = R[clamp(alpha1 + 1), beta1, alpha2, beta2]
                        R_fail_1 = R[alpha1, clamp(beta1 + 1), alpha2, beta2]
                        R1 = p1 * (1.0 + gamma * R_success_1) + (1.0 - p1) * (gamma

                        # Compute expected reward for choosing arm 2
                        p2 = alpha2 / (alpha2 + beta2)
                        R_success_2 = R[alpha1, beta1, clamp(alpha2 + 1), beta2]
                        R_fail_2 = R[alpha1, beta1, alpha2, clamp(beta2 + 1)]
                        R2 = p2 * (1.0 + gamma * R_success_2) + (1.0 - p2) * (gamma

                        # Choose the action with the higher expected reward
                        new_val = max(R1, R2)

                        # Update the value function
                        old_val = R[alpha1, beta1, alpha2, beta2]
                        diff = abs(new_val - old_val)
                        if diff > delta:
                            delta = diff
                        R[alpha1, beta1, alpha2, beta2] = new_val

                        # Update the policy
                        if R1 > R2:
                            policy[alpha1, beta1, alpha2, beta2] = 0 # Choose arm1
                        else:
                            policy[alpha1, beta1, alpha2, beta2] = 1 # Choose arm2

        if delta < tol:
            break

    return R, policy

# -----
# Simulation Parameters
# -----
true_theta = [0.3, 0.6] # True success probabilities for arm 1 and arm 2

```

```

alpha_prior = [1, 1]           # Prior alpha parameters for Beta distributions
beta_prior = [1, 1]           # Prior beta parameters for Beta distributions
gamma_values = np.linspace(0.9, 1, 100) # Gamma values
time_steps = 5000             # Number of pulls per trial
repeat_time = 10              # Number of trials per gamma
M = 17                        # Truncation level for DP

# -----
# Simulation Function for DP-Based Policy
# -----
def simulate_trial(policy, M, true_theta, alpha_prior, beta_prior, gamma, time_step
    """
    Simulate a single trial of the bandit problem using the provided policy.

    Arguments:
    -----
    policy : 4D numpy array
        Optimal policy derived from DP.
    M : int
        Truncation level.
    true_theta : list of float
        True success probabilities for each arm.
    alpha_prior : list of int
        Prior alpha parameters for Beta distributions.
    beta_prior : list of int
        Prior beta parameters for Beta distributions.
    gamma : float
        Discount factor.
    time_steps : int
        Number of pulls in the trial.

    Returns:
    -----
    total_reward : float
        Total discounted reward accumulated in the trial.
    """
    # Initialize Beta parameters
    alpha = [alpha_prior[0], alpha_prior[1]]
    beta = [beta_prior[0], beta_prior[1]]
    total_reward = 0.0
    current_gamma_power = 1.0 #  $\gamma^{t-1}$ , starts at  $t=1$ 

    for t in range(1, time_steps + 1):
        # Current state with truncation
        a1 = min(alpha[0], M)
        b1 = min(beta[0], M)
        a2 = min(alpha[1], M)
        b2 = min(beta[1], M)

        # Determine action from policy
        action = policy[a1, b1, a2, b2]

        # Pull the selected arm
        arm = action # 0 or 1
        success = np.random.rand() < true_theta[arm]
        if success:

```

```

        total_reward += current_gamma_power # Reward is  $\gamma^{t-1}$ 
        alpha[arm] += 1
    else:
        beta[arm] += 1

    # Update the discount factor for the next time step
    current_gamma_power *= gamma

    return total_reward

# -----
# Thompson Sampling Simulation Function
# -----
def thompson_sampling_simulation(true_theta, alpha_prior, beta_prior, gamma, time_steps):
    """
    Simulate a single trial of the bandit problem using Thompson Sampling.

    Arguments:
    -----
    true_theta : list of float
        True success probabilities for each arm.
    alpha_prior : list of int
        Prior alpha parameters for Beta distributions.
    beta_prior : list of int
        Prior beta parameters for Beta distributions.
    gamma : float
        Discount factor.
    time_steps : int
        Number of pulls in the trial.

    Returns:
    -----
    total_reward : float
        Total discounted reward accumulated in the trial.
    """
    alpha = np.array(alpha_prior, dtype=float)
    beta = np.array(beta_prior, dtype=float)

    total_discounted_reward = 0.0
    discount_power = 0

    for t in range(1, time_steps + 1):
        # Sample  $\hat{\theta}$  from current posterior for each arm
        sampled_thetas = np.random.beta(alpha, beta)

        # Choose the arm that maximizes the sampled  $\theta$ 
        chosen_arm = np.argmax(sampled_thetas)

        # Draw reward
        reward = 1 if np.random.rand() < true_theta[chosen_arm] else 0

        # Update posterior
        alpha[chosen_arm] += reward
        beta[chosen_arm] += (1 - reward)

        # Accumulate discounted reward

```

```

        total_discounted_reward += (gamma ** discount_power) * reward
        discount_power += 1

    return total_discounted_reward

# -----
# Gamma Evaluation Function
# -----
def evaluate_gamma(gamma):
    """
    Evaluate a single gamma value by solving DP and running simulations for both DP
    and Thompson Sampling.

    Arguments:
    -----
    gamma : float
        Discount factor.

    Returns:
    -----
    gamma : float
        The gamma value evaluated.
    average_reward_dp : float
        Average total discounted reward over all DP trials.
    average_reward_ts : float
        Average total discounted reward over all TS trials.
    """
    print(f"Evaluating gamma = {gamma:.4f}")

    # Solve DP to get the policy
    R, policy = solve_2armed_bandit_dp(M=M, gamma=gamma)

    # Initialize total rewards for all trials
    total_rewards_dp = np.zeros(repeat_time, dtype=np.float64)
    total_rewards_ts = np.zeros(repeat_time, dtype=np.float64)

    # Simulate all trials for DP-based policy
    for trial in range(1, repeat_time + 1):
        reward = simulate_trial(
            policy, M, true_theta, alpha_prior, beta_prior, gamma, time_steps
        )
        total_rewards_dp[trial - 1] = reward

    # Simulate all trials for Thompson Sampling policy
    for trial in range(1, repeat_time + 1):
        reward_ts = thompson_sampling_simulation(
            true_theta, alpha_prior, beta_prior, gamma, time_steps
        )
        total_rewards_ts[trial - 1] = reward_ts

    # Calculate average rewards
    average_reward_dp = np.mean(total_rewards_dp)
    average_reward_ts = np.mean(total_rewards_ts)

    print(f"Gamma={gamma:.4f}: DP Avg Reward={average_reward_dp:.2f}, TS Avg Reward={average_reward_ts:.2f}")

    return gamma, average_reward_dp, average_reward_ts

```

```

# -----
# Main Evaluation Loop
# -----
# Initialize lists to store results
results_dp = []
results_ts = []

# Total number of gamma values
total_gammas = len(gamma_values)

# Iterate over gamma_values and collect results
for idx, gamma in enumerate(gamma_values, 1):
    print(f"\nProcessing gamma {idx}/{total_gammas}: gamma = {gamma:.4f}")
    gamma_result = evaluate_gamma(gamma)
    _, avg_dp, avg_ts = gamma_result
    results_dp.append(avg_dp)
    results_ts.append(avg_ts)

print("\nAll gamma values have been evaluated.\n")

# Convert results to numpy arrays for easier processing
gamma_evaluated = np.array(gamma_values)
avg_rewards_dp = np.array(results_dp)
avg_rewards_ts = np.array(results_ts)

# Compute the gap between DP and TS
gap = avg_rewards_dp - avg_rewards_ts

# -----
# Find the gamma with the highest average reward for DP
# -----
optimal_index = np.argmax(avg_rewards_dp)
optimal_gamma = gamma_evaluated[optimal_index]
optimal_reward_dp = avg_rewards_dp[optimal_index]
optimal_reward_ts = avg_rewards_ts[optimal_index]

print(f"Optimal gamma for DP: {optimal_gamma:.4f}")
print(f"DP Reward at Optimal Gamma: {optimal_reward_dp:.2f}")
print(f"TS Reward at Optimal Gamma: {optimal_reward_ts:.2f}")

# -----
# Plot the Results
# -----
plt.figure(figsize=(14, 6))

# Plot Average Rewards for DP and TS
plt.subplot(1, 2, 1)
plt.plot(gamma_evaluated, avg_rewards_dp, linestyle='-', color='blue', label='DP Op
plt.plot(gamma_evaluated, avg_rewards_ts, linestyle='--', color='green', label='Tho
plt.xlabel('Gamma')
plt.ylabel('Average Discounted Reward')
plt.title('Average Discounted Reward vs Gamma')
plt.axvline(optimal_gamma, color='red', linestyle='--', label=f'Optimal Gamma: {opt
plt.legend()
plt.grid(True)

```



```
# Plot the Gap between DP and TS
plt.subplot(1, 2, 2)
plt.plot(gamma_evaluated, gap, linestyle='-', color='purple')
plt.xlabel('Gamma')
plt.ylabel('Reward Gap (DP - TS)')
plt.title('Gap Between DP Optimal Policy and Thompson Sampling')
plt.axvline(optimal_gamma, color='red', linestyle='--', label=f'Optimal Gamma: {opt}')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()
```

Processing gamma 1/100: gamma = 0.9000
Evaluating gamma = 0.9000
Gamma=0.9000: DP Avg Reward=4.23, TS Avg Reward=5.18

Processing gamma 2/100: gamma = 0.9010
Evaluating gamma = 0.9010
Gamma=0.9010: DP Avg Reward=5.87, TS Avg Reward=4.99

Processing gamma 3/100: gamma = 0.9020
Evaluating gamma = 0.9020
Gamma=0.9020: DP Avg Reward=5.06, TS Avg Reward=5.79

Processing gamma 4/100: gamma = 0.9030
Evaluating gamma = 0.9030
Gamma=0.9030: DP Avg Reward=6.04, TS Avg Reward=5.97

Processing gamma 5/100: gamma = 0.9040
Evaluating gamma = 0.9040
Gamma=0.9040: DP Avg Reward=5.80, TS Avg Reward=5.14

Processing gamma 6/100: gamma = 0.9051
Evaluating gamma = 0.9051
Gamma=0.9051: DP Avg Reward=5.22, TS Avg Reward=5.12

Processing gamma 7/100: gamma = 0.9061
Evaluating gamma = 0.9061
Gamma=0.9061: DP Avg Reward=5.56, TS Avg Reward=5.59

Processing gamma 8/100: gamma = 0.9071
Evaluating gamma = 0.9071
Gamma=0.9071: DP Avg Reward=6.20, TS Avg Reward=5.20

Processing gamma 9/100: gamma = 0.9081
Evaluating gamma = 0.9081
Gamma=0.9081: DP Avg Reward=6.30, TS Avg Reward=5.83

Processing gamma 10/100: gamma = 0.9091
Evaluating gamma = 0.9091
Gamma=0.9091: DP Avg Reward=6.32, TS Avg Reward=5.02

Processing gamma 11/100: gamma = 0.9101
Evaluating gamma = 0.9101
Gamma=0.9101: DP Avg Reward=5.29, TS Avg Reward=5.70

Processing gamma 12/100: gamma = 0.9111
Evaluating gamma = 0.9111
Gamma=0.9111: DP Avg Reward=6.39, TS Avg Reward=6.32

Processing gamma 13/100: gamma = 0.9121
Evaluating gamma = 0.9121
Gamma=0.9121: DP Avg Reward=7.09, TS Avg Reward=5.13

Processing gamma 14/100: gamma = 0.9131
Evaluating gamma = 0.9131
Gamma=0.9131: DP Avg Reward=6.68, TS Avg Reward=6.59

Processing gamma 15/100: $\gamma = 0.9141$
Evaluating gamma = 0.9141
Gamma=0.9141: DP Avg Reward=6.42, TS Avg Reward=6.76

Processing gamma 16/100: $\gamma = 0.9152$
Evaluating gamma = 0.9152
Gamma=0.9152: DP Avg Reward=6.82, TS Avg Reward=6.35

Processing gamma 17/100: $\gamma = 0.9162$
Evaluating gamma = 0.9162
Gamma=0.9162: DP Avg Reward=6.66, TS Avg Reward=5.86

Processing gamma 18/100: $\gamma = 0.9172$
Evaluating gamma = 0.9172
Gamma=0.9172: DP Avg Reward=6.40, TS Avg Reward=5.62

Processing gamma 19/100: $\gamma = 0.9182$
Evaluating gamma = 0.9182
Gamma=0.9182: DP Avg Reward=7.05, TS Avg Reward=6.29

Processing gamma 20/100: $\gamma = 0.9192$
Evaluating gamma = 0.9192
Gamma=0.9192: DP Avg Reward=5.76, TS Avg Reward=6.04

Processing gamma 21/100: $\gamma = 0.9202$
Evaluating gamma = 0.9202
Gamma=0.9202: DP Avg Reward=6.67, TS Avg Reward=6.43

Processing gamma 22/100: $\gamma = 0.9212$
Evaluating gamma = 0.9212
Gamma=0.9212: DP Avg Reward=6.55, TS Avg Reward=6.90

Processing gamma 23/100: $\gamma = 0.9222$
Evaluating gamma = 0.9222
Gamma=0.9222: DP Avg Reward=7.24, TS Avg Reward=6.80

Processing gamma 24/100: $\gamma = 0.9232$
Evaluating gamma = 0.9232
Gamma=0.9232: DP Avg Reward=7.34, TS Avg Reward=7.15

Processing gamma 25/100: $\gamma = 0.9242$
Evaluating gamma = 0.9242
Gamma=0.9242: DP Avg Reward=7.07, TS Avg Reward=6.72

Processing gamma 26/100: $\gamma = 0.9253$
Evaluating gamma = 0.9253
Gamma=0.9253: DP Avg Reward=8.15, TS Avg Reward=6.17

Processing gamma 27/100: $\gamma = 0.9263$
Evaluating gamma = 0.9263
Gamma=0.9263: DP Avg Reward=8.09, TS Avg Reward=7.71

Processing gamma 28/100: $\gamma = 0.9273$
Evaluating gamma = 0.9273
Gamma=0.9273: DP Avg Reward=8.47, TS Avg Reward=6.94

Processing gamma 29/100: gamma = 0.9283
Evaluating gamma = 0.9283
Gamma=0.9283: DP Avg Reward=7.42, TS Avg Reward=6.37

Processing gamma 30/100: gamma = 0.9293
Evaluating gamma = 0.9293
Gamma=0.9293: DP Avg Reward=7.28, TS Avg Reward=7.12

Processing gamma 31/100: gamma = 0.9303
Evaluating gamma = 0.9303
Gamma=0.9303: DP Avg Reward=7.80, TS Avg Reward=6.92

Processing gamma 32/100: gamma = 0.9313
Evaluating gamma = 0.9313
Gamma=0.9313: DP Avg Reward=7.46, TS Avg Reward=8.34

Processing gamma 33/100: gamma = 0.9323
Evaluating gamma = 0.9323
Gamma=0.9323: DP Avg Reward=8.42, TS Avg Reward=8.02

Processing gamma 34/100: gamma = 0.9333
Evaluating gamma = 0.9333
Gamma=0.9333: DP Avg Reward=7.52, TS Avg Reward=8.12

Processing gamma 35/100: gamma = 0.9343
Evaluating gamma = 0.9343
Gamma=0.9343: DP Avg Reward=8.76, TS Avg Reward=8.06

Processing gamma 36/100: gamma = 0.9354
Evaluating gamma = 0.9354
Gamma=0.9354: DP Avg Reward=8.03, TS Avg Reward=7.71

Processing gamma 37/100: gamma = 0.9364
Evaluating gamma = 0.9364
Gamma=0.9364: DP Avg Reward=8.45, TS Avg Reward=7.29

Processing gamma 38/100: gamma = 0.9374
Evaluating gamma = 0.9374
Gamma=0.9374: DP Avg Reward=9.14, TS Avg Reward=8.19

Processing gamma 39/100: gamma = 0.9384
Evaluating gamma = 0.9384
Gamma=0.9384: DP Avg Reward=10.54, TS Avg Reward=8.47

Processing gamma 40/100: gamma = 0.9394
Evaluating gamma = 0.9394
Gamma=0.9394: DP Avg Reward=8.79, TS Avg Reward=8.70

Processing gamma 41/100: gamma = 0.9404
Evaluating gamma = 0.9404
Gamma=0.9404: DP Avg Reward=8.27, TS Avg Reward=9.21

Processing gamma 42/100: gamma = 0.9414
Evaluating gamma = 0.9414
Gamma=0.9414: DP Avg Reward=9.20, TS Avg Reward=9.22

Processing gamma 43/100: gamma = 0.9424
Evaluating gamma = 0.9424
Gamma=0.9424: DP Avg Reward=9.35, TS Avg Reward=8.63

Processing gamma 44/100: gamma = 0.9434
Evaluating gamma = 0.9434
Gamma=0.9434: DP Avg Reward=9.59, TS Avg Reward=9.23

Processing gamma 45/100: gamma = 0.9444
Evaluating gamma = 0.9444
Gamma=0.9444: DP Avg Reward=9.20, TS Avg Reward=10.17

Processing gamma 46/100: gamma = 0.9455
Evaluating gamma = 0.9455
Gamma=0.9455: DP Avg Reward=10.27, TS Avg Reward=9.44

Processing gamma 47/100: gamma = 0.9465
Evaluating gamma = 0.9465
Gamma=0.9465: DP Avg Reward=9.82, TS Avg Reward=9.38

Processing gamma 48/100: gamma = 0.9475
Evaluating gamma = 0.9475
Gamma=0.9475: DP Avg Reward=9.56, TS Avg Reward=10.37

Processing gamma 49/100: gamma = 0.9485
Evaluating gamma = 0.9485
Gamma=0.9485: DP Avg Reward=9.20, TS Avg Reward=9.30

Processing gamma 50/100: gamma = 0.9495
Evaluating gamma = 0.9495
Gamma=0.9495: DP Avg Reward=10.23, TS Avg Reward=9.72

Processing gamma 51/100: gamma = 0.9505
Evaluating gamma = 0.9505
Gamma=0.9505: DP Avg Reward=11.19, TS Avg Reward=9.61

Processing gamma 52/100: gamma = 0.9515
Evaluating gamma = 0.9515
Gamma=0.9515: DP Avg Reward=10.85, TS Avg Reward=10.47

Processing gamma 53/100: gamma = 0.9525
Evaluating gamma = 0.9525
Gamma=0.9525: DP Avg Reward=11.53, TS Avg Reward=11.00

Processing gamma 54/100: gamma = 0.9535
Evaluating gamma = 0.9535
Gamma=0.9535: DP Avg Reward=12.21, TS Avg Reward=11.62

Processing gamma 55/100: gamma = 0.9545
Evaluating gamma = 0.9545
Gamma=0.9545: DP Avg Reward=9.53, TS Avg Reward=10.93

Processing gamma 56/100: gamma = 0.9556
Evaluating gamma = 0.9556
Gamma=0.9556: DP Avg Reward=11.30, TS Avg Reward=11.66

Processing gamma 57/100: $\gamma = 0.9566$
Evaluating gamma = 0.9566
Gamma=0.9566: DP Avg Reward=12.15, TS Avg Reward=12.40

Processing gamma 58/100: $\gamma = 0.9576$
Evaluating gamma = 0.9576
Gamma=0.9576: DP Avg Reward=12.47, TS Avg Reward=12.68

Processing gamma 59/100: $\gamma = 0.9586$
Evaluating gamma = 0.9586
Gamma=0.9586: DP Avg Reward=13.73, TS Avg Reward=12.77

Processing gamma 60/100: $\gamma = 0.9596$
Evaluating gamma = 0.9596
Gamma=0.9596: DP Avg Reward=14.52, TS Avg Reward=12.84

Processing gamma 61/100: $\gamma = 0.9606$
Evaluating gamma = 0.9606
Gamma=0.9606: DP Avg Reward=14.34, TS Avg Reward=13.61

Processing gamma 62/100: $\gamma = 0.9616$
Evaluating gamma = 0.9616
Gamma=0.9616: DP Avg Reward=14.11, TS Avg Reward=13.29

Processing gamma 63/100: $\gamma = 0.9626$
Evaluating gamma = 0.9626
Gamma=0.9626: DP Avg Reward=14.08, TS Avg Reward=13.75

Processing gamma 64/100: $\gamma = 0.9636$
Evaluating gamma = 0.9636
Gamma=0.9636: DP Avg Reward=15.20, TS Avg Reward=14.47

Processing gamma 65/100: $\gamma = 0.9646$
Evaluating gamma = 0.9646
Gamma=0.9646: DP Avg Reward=16.59, TS Avg Reward=15.02

Processing gamma 66/100: $\gamma = 0.9657$
Evaluating gamma = 0.9657
Gamma=0.9657: DP Avg Reward=17.22, TS Avg Reward=15.81

Processing gamma 67/100: $\gamma = 0.9667$
Evaluating gamma = 0.9667
Gamma=0.9667: DP Avg Reward=15.93, TS Avg Reward=15.92

Processing gamma 68/100: $\gamma = 0.9677$
Evaluating gamma = 0.9677
Gamma=0.9677: DP Avg Reward=18.08, TS Avg Reward=16.25

Processing gamma 69/100: $\gamma = 0.9687$
Evaluating gamma = 0.9687
Gamma=0.9687: DP Avg Reward=18.17, TS Avg Reward=17.32

Processing gamma 70/100: $\gamma = 0.9697$
Evaluating gamma = 0.9697
Gamma=0.9697: DP Avg Reward=18.94, TS Avg Reward=16.87

Processing gamma 71/100: $\gamma = 0.9707$
Evaluating gamma = 0.9707
Gamma=0.9707: DP Avg Reward=18.94, TS Avg Reward=19.12

Processing gamma 72/100: $\gamma = 0.9717$
Evaluating gamma = 0.9717
Gamma=0.9717: DP Avg Reward=21.93, TS Avg Reward=19.63

Processing gamma 73/100: $\gamma = 0.9727$
Evaluating gamma = 0.9727
Gamma=0.9727: DP Avg Reward=21.67, TS Avg Reward=19.69

Processing gamma 74/100: $\gamma = 0.9737$
Evaluating gamma = 0.9737
Gamma=0.9737: DP Avg Reward=20.85, TS Avg Reward=19.93

Processing gamma 75/100: $\gamma = 0.9747$
Evaluating gamma = 0.9747
Gamma=0.9747: DP Avg Reward=19.46, TS Avg Reward=22.04

Processing gamma 76/100: $\gamma = 0.9758$
Evaluating gamma = 0.9758
Gamma=0.9758: DP Avg Reward=25.31, TS Avg Reward=22.38

Processing gamma 77/100: $\gamma = 0.9768$
Evaluating gamma = 0.9768
Gamma=0.9768: DP Avg Reward=25.60, TS Avg Reward=22.36

Processing gamma 78/100: $\gamma = 0.9778$
Evaluating gamma = 0.9778
Gamma=0.9778: DP Avg Reward=24.24, TS Avg Reward=25.46

Processing gamma 79/100: $\gamma = 0.9788$
Evaluating gamma = 0.9788
Gamma=0.9788: DP Avg Reward=24.13, TS Avg Reward=25.34

Processing gamma 80/100: $\gamma = 0.9798$
Evaluating gamma = 0.9798
Gamma=0.9798: DP Avg Reward=28.80, TS Avg Reward=25.61

Processing gamma 81/100: $\gamma = 0.9808$
Evaluating gamma = 0.9808
Gamma=0.9808: DP Avg Reward=28.95, TS Avg Reward=27.10

Processing gamma 82/100: $\gamma = 0.9818$
Evaluating gamma = 0.9818
Gamma=0.9818: DP Avg Reward=29.40, TS Avg Reward=30.41

Processing gamma 83/100: $\gamma = 0.9828$
Evaluating gamma = 0.9828
Gamma=0.9828: DP Avg Reward=34.18, TS Avg Reward=32.91

Processing gamma 84/100: $\gamma = 0.9838$
Evaluating gamma = 0.9838
Gamma=0.9838: DP Avg Reward=34.37, TS Avg Reward=34.12

Processing gamma 85/100: $\gamma = 0.9848$
Evaluating gamma = 0.9848
Gamma=0.9848: DP Avg Reward=34.12, TS Avg Reward=37.57

Processing gamma 86/100: $\gamma = 0.9859$
Evaluating gamma = 0.9859
Gamma=0.9859: DP Avg Reward=39.85, TS Avg Reward=41.39

Processing gamma 87/100: $\gamma = 0.9869$
Evaluating gamma = 0.9869
Gamma=0.9869: DP Avg Reward=43.05, TS Avg Reward=43.11

Processing gamma 88/100: $\gamma = 0.9879$
Evaluating gamma = 0.9879
Gamma=0.9879: DP Avg Reward=46.48, TS Avg Reward=47.55

Processing gamma 89/100: $\gamma = 0.9889$
Evaluating gamma = 0.9889
Gamma=0.9889: DP Avg Reward=50.42, TS Avg Reward=51.01

Processing gamma 90/100: $\gamma = 0.9899$
Evaluating gamma = 0.9899
Gamma=0.9899: DP Avg Reward=57.54, TS Avg Reward=57.58

Processing gamma 91/100: $\gamma = 0.9909$
Evaluating gamma = 0.9909
Gamma=0.9909: DP Avg Reward=61.79, TS Avg Reward=67.40

Processing gamma 92/100: $\gamma = 0.9919$
Evaluating gamma = 0.9919
Gamma=0.9919: DP Avg Reward=70.39, TS Avg Reward=71.11

Processing gamma 93/100: $\gamma = 0.9929$
Evaluating gamma = 0.9929
Gamma=0.9929: DP Avg Reward=78.76, TS Avg Reward=79.83

Processing gamma 94/100: $\gamma = 0.9939$
Evaluating gamma = 0.9939
Gamma=0.9939: DP Avg Reward=97.94, TS Avg Reward=96.37

Processing gamma 95/100: $\gamma = 0.9949$
Evaluating gamma = 0.9949
Gamma=0.9949: DP Avg Reward=110.69, TS Avg Reward=116.20

Processing gamma 96/100: $\gamma = 0.9960$
Evaluating gamma = 0.9960
Gamma=0.9960: DP Avg Reward=133.00, TS Avg Reward=145.14

Processing gamma 97/100: $\gamma = 0.9970$
Evaluating gamma = 0.9970
Gamma=0.9970: DP Avg Reward=194.02, TS Avg Reward=193.00

Processing gamma 98/100: $\gamma = 0.9980$
Evaluating gamma = 0.9980
Gamma=0.9980: DP Avg Reward=289.31, TS Avg Reward=286.59

Processing gamma 99/100: gamma = 0.9990

Evaluating gamma = 0.9990

Gamma=0.9990: DP Avg Reward=555.59, TS Avg Reward=582.36

Processing gamma 100/100: gamma = 1.0000

Evaluating gamma = 1.0000

Gamma=1.0000: DP Avg Reward=2996.70, TS Avg Reward=3000.90

All gamma values have been evaluated.

Optimal gamma for DP: 1.0000

DP Reward at Optimal Gamma: 2996.70

TS Reward at Optimal Gamma: 3000.90

